

3.4 BANDWIDTH CONTROLS OPTIMIZATION ACCESS TO HIGH PRECISION

The construction’s nearly width-independent relative bandwidth $\lambda = \gamma h$ requires $\gamma = \Theta(W)$, yet trained slopes fall short (Figure 4). Even though larger widths admit smaller λ at comparable construction error (Figure 2b), small slopes can delay readout learning by weakening target-relevant kernel directions. We quantify this delay by replacing Section 3.1’s direct least-squares solve with GD on fixed features. This connects prior analysis of spectral bias (Xu et al., 2020; Tancik et al., 2020) to readout optimization in the quasi-interpolant dictionary.

Theorem 3.2 (Slope-dependent output-error lower bound). *For the fixed-center construction of Section 3.1 with common slope $\gamma > 0$, let $\mu_1 \geq \mu_2 \geq \dots \geq 0$ be the normalized readout-kernel eigenvalues. For a nonzero sampled target, readout GD from zero with $\eta = 1/(2\mu_1)$ satisfies*

$$\frac{\|f - q_n\|_{2,Z}^2}{\|f\|_{2,Z}^2} \geq \sum_i p_i(\gamma) \left(1 - \frac{\bar{p}_i(\gamma)}{2}\right)^{2n}, \quad n \geq 0. \quad (1)$$

Here $\|\cdot\|_{2,Z}$ is the training-sample RMS norm, p_i is the fraction of sampled target energy projected onto kernel eigenvector i , and $\mu_i/\mu_1 \leq \bar{p}_i(\gamma) \leq 1$. The gamma-dependent upper bounds $\bar{p}_i$ come from the spectral Theorem A.1 in Appendix A.

Proof sketch. The readout kernel contains normalized inner products of feature vectors at pairs of training points. To bound its spectrum, we separate the constant sample direction and approximate the sum over centers by an integral. This approximation depends explicitly on $d \coth(\gamma d)$, where d separates those points. Bounding the finite-center corrections gives $\bar{p}_i(\gamma)$; standard GD eigenmode decay then yields the output-error bound.




Figure 3: **Small slopes weaken kernel directions and delay readout learning.** Width $W = 512$; target $f(x) = [\sin(2\pi x) + \frac{1}{2}\sin(6\pi x) + \frac{1}{4}\sin(14\pi x)]/\sqrt{21/32}$ on $[-1, 1]$. **Left:** Kernel ratios μ_i/μ_1 (dots) and intervals from Theorem A.1 (dashed boundaries, shading) for three ranks carrying substantial target energy. Lines connect ordered ranks; eigenvectors change with bandwidth. **Right:** Executed frozen-readout GD errors (solid, hollow markers) and square roots of Theorem 3.2’s lower bounds (dashed). The error bound combines theorem-derived rate bounds with all resolved target projections onto the frozen kernel’s eigenvectors. Errors are relative training RMS through five million updates; geometry and protocol are in Appendix B.

Figure 3 (left) shows that small bandwidths leave target-relevant eigenvalues tiny relative to μ_1 , as captured by the appendix’s spectral bounds. Since μ_1 limits the stable GD step size, these target components decay slowly even when the dictionary admits an accurate fit. Theorem 3.2 converts those rate bounds into output-error lower bounds that closely track executed GD (Figure 3, right): bandwidth controls how quickly training accesses the dictionary’s accuracy.

The acquisition shortfall also appears empirically under Adam: joint training leaves a bandwidth-acquisition gap and output error above its frozen-feature counterpart (Figure 4). This motivates the next question: why does joint training fail to acquire useful slopes and centers together?

3.5 TRAINING MUST ACQUIRE USEFUL SLOPES AND CENTERS

Joint training also leaves a substantial precision gap over this budget. After five million updates, joint Adam reaches median dense-grid relative error 1.81×10^{-3} , versus 6.62×10^{-7} on the supplied $\lambda = 1/4$ features (Figure 4, left).




Figure 4: **Partial slope acquisition leaves a precision gap.** **Left:** Joint Adam/GD relative output errors and frozen uniform $\lambda = 1/4$ references under the same five-million-update budget. **Right:** RMS and 99th-percentile scaled slopes $h[a_j]$ in the same joint runs. The reference is a supplied geometry, not a necessary slope threshold. Joint curves summarize five paired seeds; shading retains seed variation and recorded within-bin extrema. Cosine schedules span the full horizon. These are empirical comparisons; Appendix B gives the protocol.

Adam’s 99th-percentile scaled slope approaches $1/4$, but only 4–7 of 512 features reach that reference across the five seeds (Figure 4, right). Readout refits on the learned features barely improve their error (Appendix B). Longer training helps, but acquiring a few steep features does not recover the supplied geometry’s accuracy within this budget. This motivates the next analysis: training must acquire useful slopes and centers together. The construction supplies both the feature geometry and the readout directly.

A Explicit spectral bounds and proof

The main theorem concerns the actual finite tanh dictionary. An auxiliary center integral exposes its gamma dependence; lattice and exterior-center allowances transfer its spectrum to that dictionary. This appendix specifies each approximation and the resulting endpoints. The scale-dependent spectral bias is related to [Xu et al. \(2020\)](#); [Tancik et al. \(2020\)](#); the argument below quantifies it for the finite quasi-interpolant dictionary.

Finite setup. Let $z_1, \dots, z_M$ be distinct training points with $M \geq 2$, W consecutive centers on $c_0 + h\mathbb{Z}$ cover their interval, and $\gamma > 0$. Define

$$\Phi_\gamma = [\mathbf{1}, (\tanh(\gamma(z_a - x_j)))_{a,j}], \quad K_\gamma = \frac{1}{M} \Phi_\gamma \Phi_\gamma^\top, \quad y = (f(z_a))_{a=1}^M \neq 0. \quad (2)$$

Readout GD minimizes $\|\Phi_\gamma w - y\|^2 / (2M)$. Write its descending kernel eigenpairs as (μ_i, u_i) and define

$$\rho_i = \mu_i / \mu_1, \quad p_i = |u_i^\top y|^2 / \|y\|^2, \quad e_n^2 = \frac{\|\Phi_\gamma w_n - y\|^2}{\|y\|^2} = \frac{\|f - q_n\|_{2,Z}^2}{\|f\|_{2,Z}^2}. \quad (3)$$

The energy weights sum to one, including any exact nullspace component. The theorem permits arbitrary distinct samples; the reported experiments use uniform samples. Write $v_0 = \mathbf{1} / \sqrt{M}$ and let Q have orthonormal columns spanning $v_0^\perp$. This projection removes the common constant sample pattern, not any target assumption. Put $t_\gamma(c) = (\tanh(\gamma(z_a - c)))_a$.

The gamma-dependent correlation model. Its center-integrated, mean-zero kernel is

$$H_\gamma = \frac{1}{hM} \int_{\mathbb{R}} Q^\top t_\gamma(c) t_\gamma(c)^\top Q dc = -\frac{2}{hM} Q^\top [d_{ab} \coth(\gamma d_{ab})]_{a,b} Q, \quad d_{ab} = z_a - z_b. \quad (4)$$

The diagonal of the bracketed matrix is $1/\gamma$. Choose a finite set $\mathcal{F}$ of exterior lattice centers absent from the actual dictionary, leaving two unretained tails starting at $c_L < \min z_a$ and $c_R > \max z_a$. Subtract their retained contribution:

$$T_\gamma = \frac{1}{M} \sum_{c \in \mathcal{F}} Q^\top t_\gamma(c) t_\gamma(c)^\top Q, \quad S_\gamma = H_\gamma - T_\gamma. \quad (5)$$

Let $\beta_1 \geq \dots \geq \beta_{M-1}$ be the eigenvalues of S_γ . For any $0 < \theta < \pi/2$, the lattice and remaining-tail allowances are

$$\delta_\gamma = \frac{8\gamma \|z - \bar{z}\mathbf{1}\|^2 \sec^4 \theta}{3hM [\exp(2\pi\theta/(\gamma h)) - 1]}, \quad \bar{z} = M^{-1} \sum_a z_a, \quad (6)$$

$$\tau_\gamma = \frac{4[e^{-4\gamma(c_R - z_{\max})} + e^{-4\gamma(z_{\min} - c_L)}]}{1 - e^{-4\gamma h}}. \quad (7)$$

They give the matrix enclosure

$$S_\gamma - (\delta_\gamma + \tau_\gamma)I \preceq Q^\top K_\gamma Q \preceq S_\gamma + \delta_\gamma I. \quad (8)$$

Relative-eigenvalue endpoints. Normalize using scalar bounds $\ell_\gamma \leq \mu_1 \leq L_\gamma$. The following require only matrix-vector products with the finite dictionary:

$$\ell_\gamma = v_0^\top K_\gamma v_0 \geq 1, \quad b_\gamma = \|Q^\top K_\gamma v_0\|, \quad c_\gamma = \beta_1 + \delta_\gamma, \quad (9)$$

$$L_\gamma = \frac{\ell_\gamma + c_\gamma + \sqrt{(\ell_\gamma - c_\gamma)^2 + 4b_\gamma^2}}{2}. \quad (10)$$

Set $\rho_1 = \bar{\rho}_1 = 1$ and $\beta_M = 0$.

Theorem A.1 (Slope-dependent relative-eigenvalue bounds).

$$\boxed{\rho_i = \frac{[\beta_i - \delta_\gamma - \tau_\gamma]_+}{L_\gamma}, \quad \bar{\rho}_i = \min\left\{1, \frac{\beta_{i-1} + \delta_\gamma}{\ell_\gamma}\right\}, \quad 2 \leq i \leq M.} \quad (11)$$

For the finite setup above, these endpoints enclose every kernel ratio: $\rho_i \leq \mu_i/\mu_1 \leq \bar{\rho}_i$. Exact zeros known from $\text{rank } K_\gamma \leq W + 1$ may have both endpoints set to zero.

No target or optimization trajectory enters these spectral bounds. The upper endpoints supply Theorem 3.2's output-error lower bound.

A.1 DERIVATION OF THE EXPLICIT GAMMA DEPENDENCE

The identity $1 - \tanh u \tanh v = \coth(u - v)(\tanh u - \tanh v)$ gives

$$\int_{\mathbb{R}} [1 - \tanh(\gamma(z_a - c)) \tanh(\gamma(z_b - c))] dc = 2d_{ab} \coth(\gamma d_{ab}). \quad (12)$$

Multiplication by $Q^\top$ and Q removes the constant term and proves (4). The integral is finite after projection because $Q^\top t_\gamma(c)$ decays at both ends.

The smoothing interpretation can also be made exact. With $s_\gamma(x) = \frac{\gamma}{2} \text{sech}^2(\gamma x)$, $\tanh(\gamma \cdot) = s_\gamma * \text{sign}$. Its Fourier multiplier is

$$A_\gamma(\omega) = \hat{s}_\gamma(\omega) = \frac{\pi|\omega|/(2\gamma)}{\sinh(\pi|\omega|/(2\gamma))}, \quad A_\gamma(0) = 1. \quad (13)$$

For completeness, substitution $r = e^{2\gamma x}$ in the Fourier transform gives $\int_0^\infty r^{-i\omega/(2\gamma)}(1+r)^{-2} dr$. The beta integral evaluates this as $\Gamma(1 - i\omega/(2\gamma))\Gamma(1 + i\omega/(2\gamma))$, which equals (13) by the reflection formula. For $u = Qv$, the step combination $\sum_a u_a \text{sign}(c - z_a)$ has compact support, and its derivative is $2 \sum_a u_a \delta_{z_a}$. Parseval gives

$$v^\top H_\gamma v = \frac{2}{\pi h M} \int_{\mathbb{R}} \frac{A_\gamma(\omega)^2}{\omega^2} \left| \sum_a (Qv)_a e^{-i\omega z_a} \right|^2 d\omega. \quad (14)$$

The apparent zero-frequency singularity is removable since $\sum_a (Qv)_a = 0$. Thus gamma acts on a direction's frequency content through the explicit weight A_γ^2 . Since $t/\sinh t$ decreases for $t > 0$, increasing gamma increases every directional response of H_γ and hence its ordered eigenvalues. This order concerns the integral kernel; exterior subtraction and normalization need not preserve monotonicity of every finite ratio.

A.2 TRANSFER TO THE FINITE DICTIONARY

For $u \perp \mathbf{1}$, the function $g(\zeta) = \sum_a u_a \tanh(\gamma(\zeta - z_a))$ is analytic for $|\text{Im } \zeta| < \pi/(2\gamma)$. Subtracting the common translate at $\bar{z}$ and integrating derivatives, Minkowski and Cauchy-Schwarz yield

$$\int_{\mathbb{R}} |g(t \pm i\theta/\gamma)|^2 dt \leq \frac{4\gamma}{3} \sec^4 \theta \|z - \bar{z}\mathbf{1}\|^2 \|u\|^2. \quad (15)$$

Here $|\text{sech}^2(t + i\theta)| \leq \sec^2 \theta \text{sech}^2 t$ and $\int \text{sech}^4 t dt = 4/3$. Shift the contour for the analytic function $g(\zeta)^2$, not $|g(\zeta)|^2$: its transform is bounded by the right side of (15) times $e^{-\theta|\omega|/\gamma}$. Poisson summation at frequencies $2\pi k/h$ then gives

$$\left\| \frac{1}{M} \sum_{c \in c_0 + h\mathbb{Z}} Q^\top t_\gamma(c) t_\gamma(c)^\top Q - H_\gamma \right\| \leq \delta_\gamma. \quad (16)$$

The geometric sum over $k \neq 0$ produces the denominator in (6).

Deleting exterior centers converts the infinite lattice to the actual dictionary. The retained deletions give T_γ . For a remaining center $c \geq c_R$, projection removes the constant limit and $1 - \tanh t \leq 2e^{-2t}$ gives an outer-product norm, divided by M , of at most $4e^{-4\gamma(c-z_{\max})}$. Summing this geometric series and its left-tail counterpart bounds the remaining positive deletion by $\tau_\gamma I$. This proves (8).

If κ_i are the eigenvalues of $Q^\top K_\gamma Q$, min-max and codimension-one interlacing give

$$\beta_i - \delta_\gamma - \tau_\gamma \leq \kappa_i \leq \beta_i + \delta_\gamma, \quad \mu_i \geq \kappa_i \geq \mu_{i+1}. \quad (17)$$

The Rayleigh quotient gives $\mu_1 \geq \ell_\gamma$. In the decomposition $\text{span}(v_0) \oplus v_0^\perp$, the diagonal blocks of K_γ are bounded above by ℓ_γ and $c_\gamma I$, and the off-diagonal block has norm b_γ . Its quadratic form is therefore bounded by the largest eigenvalue of $\begin{bmatrix} \ell_\gamma & b_\gamma \\ b_\gamma & c_\gamma \end{bmatrix}$, namely L_γ . Dividing the interlacing bounds by these normalization endpoints proves the relative-rate enclosure from (11).

A.3 CLASSICAL OUTPUT-ERROR CONSEQUENCE

For $r_n = \Phi_\gamma w_n - y$, normalized least-squares GD gives $r_{n+1} = (I - \eta K_\gamma) r_n$ and $r_0 = -y$. Expansion in the actual finite-kernel eigenvectors gives

$$e_n^2 = \sum_i p_i (1 - \eta \mu_i)^{2n} = \sum_i p_i (1 - \rho_i/2)^{2n} \geq \sum_i p_i (1 - \bar{\rho}_i/2)^{2n}. \quad (18)$$

The last inequality follows because the factor decreases with $\rho_i \in [0, 1]$. This proves Theorem 3.2. Our step is a conservative fraction of the stability limit $2/\mu_1$; any nonoscillatory step $0 < \eta \mu_1 \leq 1$ instead uses $\eta \mu_1 \rho_i$. Only exact nullspace energy gives a permanent floor. For $\mu_i > 0$, $u_i = \Phi_\gamma(\Phi_\gamma^\top u_i / (M \mu_i))$ is representable, however small its rate. The target weights require the actual eigenvectors; the theorem bounds eigenvalues, not their target alignment.

B Evaluation, accuracy, and empirical scope

Which quantities are computed. The theorem endpoints use the explicit center integral, exterior-center subtraction, and scalar normalization bounds. A separate SVD of the actual finite feature matrix supplies the validation eigenvalues and target weights p_i . The plotted lower error uses the theorem's upper rates, not actual rates. In Figure 3 (left), ranks 4, 14, 30 carry approximately 46–48%, 8–14%, and 2.5–4.4% of sampled target energy, respectively, across the four evaluated bandwidths. The error bound uses all numerically resolved target projections.

The integral is evaluated through a positive quadrature factor rather than by subtracting nearly equal entries of (4). There is no Fourier truncation. Truncating the center integral at $\pm(\max_a |z_a| + P/\gamma)$ omits a positive matrix of norm at most $2e^{-4P}/(h\gamma)$, added to the numerical upper endpoint. The exterior remainder enters the lower endpoint. Compressing the combined integral and exterior factors Z to rank r introduces the two-sided bound $2\sigma_1(Z)\sigma_{r+1}(Z) + \sigma_{r+1}(Z)^2$ on their signed Gram matrix. Quadrature refinement and an empirical arithmetic guard are checked separately; the figures are FP64 evaluations, not interval-arithmetic certificates. Refining quadrature order from 10 to 16 and padding from 20 to 24 changes the lower error curves by at most 9.6×10^{-11} .

The remaining analytic slack comes from the interlacing index shift, the normalization bounds ℓ_γ, L_γ , and additive allowances. Small absolute errors can dominate extremely small eigenvalues; the enclosure is valid but need not be informative for every geometry or mode. In this experiment it is informative both for the plotted ratios and for the full error trajectory. Omitting unresolved target energy from the lower curve only weakens it; it is not interpreted as an exact approximation floor.

Target, geometry, and measurements. All displayed data use

$$f(x) = \frac{\sin(2\pi x) + \frac{1}{2}\sin(6\pi x) + \frac{1}{4}\sin(14\pi x)}{\sqrt{21/32}}, \quad x \in [-1, 1]. \quad (19)$$

The total hidden width is $W = 512$, comprising $N = 467$ interior cells and $R = 22$ halo centers on each side, so $h = 2/467$ and $W = N + 2R + 1$. There are 2,048 uniform training midpoints, 4,096 validation points, and 8,192 dense evaluation points. Every reported error is relative output L^2 error on its stated grid, not squared loss or an EMA. The dense grid checks resolution; it is not an untouched generalization test. Frozen GD starts its readout at zero and executes five million full-batch FP64 updates with $\eta = 1/(2\mu_1)$ at each bandwidth. Figure 3 (right) displays training error. Its spectral reference agrees with executed GD within 2.9×10^{-14} .

Joint training optimizes all slopes, centers (through affine biases), and readout parameters for five million updates on five paired seeds. One recipe per optimizer is selected by median final validation error over all five seeds. The selected Adam and GD schedules are full-horizon cosine decays with initial learning rates 0.05 and 0.2, respectively. Frozen Adam uses the same final-validation selection criterion. Figure 4 (left) shows raw training errors summarized for display; quoted joint and frozen-Adam endpoint values use the dense grid. Display bins retain extrema, including spikes, rather than replacing the comparison with first tolerance crossings. Figure 4 (right) measures $h|a_j|$ for learned features $\tanh(a_j x + b_j)$, using the original reference spacing h .

What the joint experiment establishes. Adam’s median scaled RMS slope is 0.0582, and its 99th percentile is 0.2466; the five seeds have 6, 5, 5, 4, 7 features at or above $1/4$. These statistics describe heterogeneous geometry, not an effective uniform bandwidth. The uniform $1/4$ dictionary is an empirical comparison, not a necessary scale for every target or learned dictionary.

Joint Adam continues improving late in training; its median endpoint error falls about 28% between the two- and five-million-update experiments. The result is a precision gap under a large finite budget, not permanent failure. Checked direct fits and an additional two-million-update readout restart on the learned five-million-update features barely improve their error. Multiplying the learned slopes by 4 or 16 at fixed centers worsens the median refitted error to 0.0916 or 0.139, respectively. These interventions change nonnested feature spaces: they support investigating slopes and centers together, not a theorem that increasing slopes always helps. Neither the fixed-kernel proof nor its scalar GD rates apply to evolving features or Adam’s adaptive preconditioner. The next theoretical question is how evolving features acquire a useful geometry.

Reproducibility. Figure sources, numerical provenance, protocols, and checks are indexed in `output/diagnostics/section34_long_horizon/README.md`. The figures reuse completed runs.

References

- Matthew Tancik, Pratul Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems*, volume 33, 2020. URL https://papers.neurips.cc/paper_files/paper/2020/hash/55053683268957697aa39fba6f231c68-Abstract.html.
- Zhi-Qin John Xu, Yaoyu Zhang, Tao Luo, Yanyang Xiao, and Zheng Ma. Frequency principle: Fourier analysis sheds light on deep neural networks. *Communications in Computational Physics*, 28(5):1746–1767, 2020. doi: 10.4208/cicp.OA-2020-0085. URL <https://www.global-sci.com/cicp/article/view/6896>.