

3.4 BANDWIDTH CONTROLS OPTIMIZATION ACCESS TO HIGH PRECISION

The construction’s nearly width-independent relative bandwidth $\lambda = \gamma h$ requires $\gamma = \Theta(W)$. Even though larger widths admit smaller λ at comparable construction error (Figure 2b), small slopes can delay readout optimization. We quantify this in the dictionary of Section 3.1, replacing the direct normalized least-squares solve with GD while keeping the features fixed.

Theorem 3.2 (Slope-dependent output-error lower bound). *For the fixed-center construction of Section 3.1 with common slope $\gamma > 0$ and nonzero sampled target, readout GD from zero with $\eta = 1/(2\mu_1)$ satisfies*

$$\frac{\|f - q_n\|_{2,Z}^2}{\|f\|_{2,Z}^2} \geq \sum_i p_i(\gamma) \left(1 - \frac{\bar{\rho}_i(\gamma)}{2}\right)^{2n}, \quad n \geq 0. \quad (1)$$

Here $\|\cdot\|_{2,Z}$ is the training-sample RMS norm, μ_i are the descending readout-kernel eigenvalues, p_i is the fraction of target energy in eigenvector i , and $\bar{\rho}_i(\gamma) \in [0, 1]$ bounds μ_i/μ_1 from above. The explicit gamma-dependent rate bounds and finite-geometry allowances are given in Appendix A.

Proof sketch. Gamma enters the feature correlations through $d \coth(\gamma d)$, where d is the separation between training points. Controlling the finite-center corrections gives the relative-rate bounds; a standard spectral argument then yields the output-error inequality.




Figure 3: **Small slopes weaken kernel directions and delay readout learning.** **A:** finite-kernel ratios (dots) and theorem intervals (dashed boundaries, shading) at four bandwidths. Lines connect ordered ranks; eigenvectors change with bandwidth. **B:** executed frozen-GD relative output errors (solid) and square roots of the theorem’s lower bounds (dashed), through five million updates. Both panels use width 512 and the same target and geometry (Appendix B).

Increasing λ from 1/32 to 1/4 raises the thirtieth relative eigenvalue by 1.4×10^5 , versus only 1.3 for the fourth (Figure 3A); the thirtieth mode carries about 2.5% of target energy at both endpoints.

At $\lambda = 1/32$, a direct readout fit attains 6.2×10^{-10} dense-grid relative error, yet five million GD updates retain training error 0.210 against the theorem’s lower bound 0.2098 (Figure 3B). At $\lambda = 1/4$, GD reaches 3.88×10^{-4} over the same budget. Across the four bandwidths, the bound remains at least 75% of executed error throughout training: small slopes delay accuracy already available in the dictionary.

3.5 TRAINING MUST ACQUIRE USEFUL SLOPES AND CENTERS

Joint training also leaves a substantial precision gap over this budget. After five million updates, joint Adam reaches median dense-grid relative error 1.81×10^{-3} , versus 6.62×10^{-7} on the supplied $\lambda = 1/4$ features (Figure 4A).




Figure 4: **Partial slope acquisition leaves a precision gap.** **A:** joint Adam/GD relative output errors and frozen uniform $\lambda = 1/4$ references under the same five-million-update budget. **B:** RMS and 99th-percentile scaled slopes $h|a_j|$ in A’s joint runs. The reference is a supplied geometry, not a necessary slope threshold. Joint curves summarize five paired seeds; shading retains seed variation and recorded within-bin extrema. Cosine schedules span the full horizon. These are empirical comparisons; Appendix B gives the protocol.

Adam’s 99th-percentile scaled slope approaches $1/4$, but only 4–7 of 512 features reach that reference across the five seeds (Figure 4B). Readout refits on the learned features barely improve their error (Appendix B). Longer training helps, but acquiring a few steep features does not recover the supplied geometry’s accuracy within this budget. This motivates the next analysis: training must acquire useful slopes and centers together. The construction supplies both the feature geometry and the readout directly.