

3.4 BANDWIDTH CONTROLS OPTIMIZATION ACCESS TO HIGH PRECISION

The construction’s nearly width-independent relative bandwidth $\lambda = \gamma h$ requires $\gamma = \Theta(W)$, yet trained slopes fall short (Figure 4). Even though larger widths admit smaller λ at comparable construction error (Figure 2b), small slopes can delay readout learning by weakening target-relevant kernel directions. We quantify this delay by replacing Section 3.1’s direct least-squares solve with GD on fixed features. This connects prior analysis of spectral bias (Xu et al., 2020; Tancik et al., 2020) to readout optimization in the quasi-interpolant dictionary.

Theorem 3.2 (Slope-dependent output-error lower bound). *For the fixed-center construction of Section 3.1 with common slope $\gamma > 0$, let $\mu_1 \geq \mu_2 \geq \dots \geq 0$ be the normalized readout-kernel eigenvalues. For a nonzero sampled target, readout GD from zero with $\eta = 1/(2\mu_1)$ satisfies*

$$\frac{\|f - q_n\|_{2,Z}^2}{\|f\|_{2,Z}^2} \geq \sum_i p_i(\gamma) \left(1 - \frac{\bar{p}_i(\gamma)}{2}\right)^{2n}, \quad n \geq 0. \quad (1)$$

Here $\|\cdot\|_{2,Z}$ is the training-sample RMS norm, p_i is the fraction of sampled target energy projected onto kernel eigenvector i , and $\mu_i/\mu_1 \leq \bar{p}_i(\gamma) \leq 1$. The gamma-dependent upper bounds $\bar{p}_i$ come from the spectral Theorem A.1 in Appendix A.

Proof sketch. The readout kernel contains normalized inner products of feature vectors at pairs of training points. To bound its spectrum, we separate the constant sample direction and approximate the sum over centers by an integral. This approximation depends explicitly on $d \coth(\gamma d)$, where d separates those points. Bounding the finite-center corrections gives $\bar{p}_i(\gamma)$; standard GD eigenmode decay then yields the output-error bound.




Figure 3: **Small slopes weaken kernel directions and delay readout learning.** Width $W = 512$; target $f(x) = [\sin(2\pi x) + \frac{1}{2}\sin(6\pi x) + \frac{1}{4}\sin(14\pi x)]/\sqrt{21/32}$ on $[-1, 1]$. **Left:** Kernel ratios μ_i/μ_1 (dots) and intervals from Theorem A.1 (dashed boundaries, shading) for three ranks carrying substantial target energy. Lines connect ordered ranks; eigenvectors change with bandwidth. **Right:** Executed frozen-readout GD errors (solid, hollow markers) and square roots of Theorem 3.2’s lower bounds (dashed). The error bound combines theorem-derived rate bounds with all resolved target projections onto the frozen kernel’s eigenvectors. Errors are relative training RMS through five million updates; geometry and protocol are in Appendix B.

Figure 3 (left) shows that small bandwidths leave target-relevant eigenvalues tiny relative to μ_1 , as captured by the appendix’s spectral bounds. Since μ_1 limits the stable GD step size, these target components decay slowly even when the dictionary admits an accurate fit. Theorem 3.2 converts those rate bounds into output-error lower bounds that closely track executed GD (Figure 3, right): bandwidth controls how quickly training accesses the dictionary’s accuracy.

The acquisition shortfall also appears empirically under Adam: joint training leaves a bandwidth-acquisition gap and output error above its frozen-feature counterpart (Figure 4). This motivates the next question: why does joint training fail to acquire useful slopes and centers together?

3.5 TRAINING MUST ACQUIRE USEFUL SLOPES AND CENTERS

Joint training also leaves a substantial precision gap over this budget. After five million updates, joint Adam reaches median dense-grid relative error 1.81×10^{-3} , versus 6.62×10^{-7} on the supplied $\lambda = 1/4$ features (Figure 4, left).




Figure 4: **Partial slope acquisition leaves a precision gap.** **Left:** Joint Adam/GD relative output errors and frozen uniform $\lambda = 1/4$ references under the same five-million-update budget. **Right:** RMS and 99th-percentile scaled slopes $h[a_j]$ in the same joint runs. The reference is a supplied geometry, not a necessary slope threshold. Joint curves summarize five paired seeds; shading retains seed variation and recorded within-bin extrema. Cosine schedules span the full horizon. These are empirical comparisons; Appendix B gives the protocol.

Adam’s 99th-percentile scaled slope approaches $1/4$, but only 4–7 of 512 features reach that reference across the five seeds (Figure 4, right). Readout refits on the learned features barely improve their error (Appendix B). Longer training helps, but acquiring a few steep features does not recover the supplied geometry’s accuracy within this budget. This motivates the next analysis: training must acquire useful slopes and centers together. The construction supplies both the feature geometry and the readout directly.