

3.5 WHY JOINT TRAINING ACQUIRES USEFUL FEATURE SCALES SLOWLY

Over five million updates, trained mean slopes do not keep pace with the construction scale as width grows (Figure 4a). On mixed sine at $W = 512$, joint Adam reaches median relative training error 1.81×10^{-3} , compared with 6.62×10^{-7} using the supplied $\lambda = 1/4$ dictionary; joint GD remains at 0.244 (Figure 4b). Numerically resolved readout refits barely improve the learned Adam dictionary, and increasing its slopes at fixed centers does not close the gap (Appendix B). These interventions indicate that useful slopes and centers must be acquired together.




Figure 4: **Joint training leaves a scale and accuracy gap.** Panels (a,b): mixed-sine target on $[-1, 1]$; five million updates. Joint training uses five seeds and one recipe per width and optimizer, selected by median final validation error. **(a)** Mean absolute slope versus total width, including halos. Faint curves show seed means; outlined markers show pooled means. **(b)** At $W = 512$, relative training RMS error for joint training (solid) and frozen readouts on the supplied $\lambda = 1/4$ dictionary (dashed). Joint curves show seed medians after binning. Shading preserves temporal extrema and, for joint training, seed spread. **(c)** Adam intervention responses over 10k updates from a shared 130k checkpoint. Each point compares an intervention with its matched native control: accumulated fine-slope path versus percentage difference in final RMS slope. The 48 comparisons span six target families, widths 705 and 1409, and two seeds; they are separate from the experiments in (a,b). Greater motion can increase acquisition but need not produce a substantially larger endpoint scale. All altered pulse-end RMS bandwidths remain below 0.013, compared with the construction reference 0.25. The dotted line in (a) marks that supplied construction scale, not a universal acquisition threshold. Protocol and target for (a,b) are in Appendix B.

We quantify one mechanism that limits this acquisition. In the small-parameter regime, $\tanh$'s leading non-affine term is cubic, so its target-driven growth depends on the target's low-degree polynomial projection. Bounding this coupling, higher-order contributions, and population concentration controls the growth of slopes, biases, and readouts. The resulting bound on non-affine output yields an output-error lower bound.

Theorem 3.3 (Output error under slow feature acquisition). *Let $f \neq 0$ be the sampled target and $f_{\perp}$ its component orthogonal to affine functions. Under the population-concentration, target-coupling, and coarse-tracking bounds of Appendix A, joint GD satisfies*

$$\frac{\|f - q_n\|_{2,Z}}{\|f\|_{2,Z}} \geq \frac{[\|f_{\perp}\|_{2,Z} - Q_n(B_n)]_+}{\|f\|_{2,Z}}. \quad (1)$$

Here B_n is an explicit scalar upper bound on the joint norm of hidden slopes, biases, and readouts, and $Q_n(B_n)$ bounds their non-affine output. The same comparison gives $\bar{\lambda}_n \leq h B_n / \sqrt{W}$, with reference spacing h and $\bar{\lambda}_n = W^{-1} \sum_j h |\gamma_{j,n}|$. The recurrence, coefficients, and precise conditions are given in the appendix.




Figure 5: **Limited population growth leaves a substantial output-error floor. Left:** Joint parameter norm and its upper bound B_n , normalized by the initial norm. **Right:** Executed GD relative training RMS error and the theorem’s lower bound. Solid curves and markers show GD; dashed curves show bounds. Archived $W = 705$, seed-30 continuations start after 20,000 updates and run 100,000 more at step size 0.002. Mixed sine is the target used in Figure 4; smooth step has the largest growth allowance among the six principal targets. Bounds use structural coefficients evaluated along the saved trajectories.

The conditional bounds closely track the separate GD continuation experiments: across ten continuations spanning six target families, the output-error lower bound captures at least 97.7% of executed error at every stored diagnostic time (Figure 5; Appendix B). The structural coefficients are evaluated along these trajectories. The longer Adam experiments establish a corresponding empirical acquisition gap; the theorem concerns GD.

QUILL supplies the feature geometry and readout directly, avoiding this acquisition stage. We next compose these constructed primitives into arithmetic circuits.

A Joint feature acquisition: conditions and full proof

The result controls an evolving tanh network in the ordinary Euclidean GD metric. A scalar comparison bounds the collective size of slopes, hidden biases, and readouts; a global remainder inequality converts that bound into output error. Its coefficients can be supplied as structural upper bounds or evaluated along a trajectory. The latter use is a conditional explanation, not a forecast from the initial parameters alone.

A.1 NETWORK, MEASURE, AND COARSE COMPENSATION

Let ν be a probability measure on $[-1, 1]$ with positive input variance, and write $\|\cdot\| = \|\cdot\|_{L^2(\nu)}$ throughout this appendix. For the training measure, this is the sample RMS norm $\|\cdot\|_{2,Z}$. Consider

$$\begin{aligned} q_\theta(x) &= d + \sum_{j=1}^W c_j \tanh(\gamma_j x + b_j), & L(\theta) &= \frac{1}{2} \|q_\theta - f\|^2, \\ \theta_{n+1} &= \theta_n - \eta_n \nabla L(\theta_n), & \eta_n &\geq 0. \end{aligned} \quad (\text{A.1})$$

The target lies in $L^2(\nu)$ with $\|f\| > 0$, and $[a]_+ = \max(a, 0)$. Let P_C project onto affine functions, $P_H = I - P_C$, $f_\perp = P_H f$, $q_H = P_H q_\theta$, and $e_C = P_C(q_\theta - f)$, $e_H = q_H - f_\perp$. Jacobians $J_C = D_\theta(P_C q_\theta)$ and $J_H = D_\theta q_H$ map the Euclidean parameter space into these output spaces. Their adjoints use the stated output measure; empirical adjoints therefore include the sample normalization. At the iterates under consideration assume $K_C = J_C J_C^*$ is invertible on the two-dimensional affine space. Define

$$\begin{aligned} \Pi &= I - J_C^* K_C^{-1} J_C, & \ell &= K_C^{-1} J_C J_H^* e_H, \\ F &= \Pi J_H^* e_H, & R &= J_C^* (e_C + \ell). \end{aligned} \quad (\text{A.2})$$

Then $\nabla L = F + R$, Π is an orthogonal projection, and $J_C F = 0$. Thus F includes the compensation needed to preserve affine output to first order. Small tracking R does not mean that compensation vanishes, or that e_C is exactly zero. The proof uses $\|\Pi\| \leq 1$; it needs no quantitative lower singular-value bound on J_C . At a singular coarse Gram matrix this formulation is not invoked.

Write $\vartheta = (\gamma, b, c)$ for the hidden parameter blocks, omitting the output bias, and set

$$r_j = (\gamma_j^2 + b_j^2 + c_j^2)^{1/2}, \quad S_k = \sum_j r_j^k, \quad s = \|\vartheta\|_2 = \sqrt{S_2} > 0. \quad (\text{A.3})$$

The slope and center parameters are free to evolve. Here b_j is the hidden intercept; a center, when $\gamma_j \neq 0$, is $-b_j/\gamma_j$. No bound on center displacement is inferred by dividing by a small slope.

A.2 GLOBAL TANH AND DERIVATIVE REMAINDER BOUNDS

Put $u_j = \gamma_j x + b_j$ and define

$$\Psi_3 = -\frac{1}{3} \sum_j c_j u_j^3, \quad \Psi_5 = \frac{2}{15} \sum_j c_j u_j^5, \quad J_k = D_\theta(P_H \Psi_k), \quad G_k = J_k^* f_\perp \quad (k = 3, 5). \quad (\text{A.4})$$

Set $t_3 = 1/3$, $t_5 = 2/15$, $t_7 = 17/315$ and

$$A_k = t_k \frac{2^{k/2} k^{k/2}}{(k+1)^{(k+1)/2}}, \quad C_k = (k+1)A_k, \quad k = 3, 5, 7. \quad (\text{A.5})$$

In particular $A_3 = \sqrt{6}/8$. For every real u , the remainders after the linear, cubic, and quintic tanh polynomials satisfy

$$|\mathcal{R}_k(u)| \leq t_k |u|^k, \quad |\mathcal{R}'_k(u)| \leq k t_k |u|^{k-1}, \quad k = 3, 5, 7. \quad (\text{A.6})$$

Here $\mathcal{R}_3 = \tanh u - u$, $\mathcal{R}_5 = \tanh u - u + u^3/3$, and $\mathcal{R}_7 = \tanh u - u + u^3/3 - 2u^5/15$.

Proof of the scalar bounds. Since $|\tanh u| \leq |u|$, $|\mathcal{R}'_3| = \tanh^2 u \leq u^2$; integration gives the cubic bound. Next, $\mathcal{R}'_5 = u^2 - \tanh^2 u$, whose magnitude is at most $|u - \tanh u| |u + \tanh u| \leq 2|u|^4/3$; integrate again. For the last pair, $\sup_{u \in \mathbb{R}} |\tanh^{(7)} u| \leq 272$. To check this constant explicitly, write $z = \tanh^2 u \in [0, 1]$:

$$\tanh^{(7)} u = -272 + 3968z - 12096z^2 + 13440z^3 - 5040z^4.$$

The degree-four Bernstein coefficients on $[0, 1/2]$ and $[1/2, 1]$ are $(-272, 224, 216, 124, 53)$ and $(53, -18, -68, 8, 0)$, respectively. The Bernstein basis is nonnegative and sums to one, proving the bound. Taylor's integral remainder for $\tanh$ through degree six and its derivative through degree five gives $272/7! = 17/315$ and $272/6! = 17/45$. The intervening coefficients vanish by parity.

From neurons to output and Jacobians. On $[-1, 1]$, $|u_j| \leq \sqrt{2}(\gamma_j^2 + b_j^2)^{1/2}$. At fixed r_j , maximizing $t_k 2^{k/2} |c_j| (\gamma_j^2 + b_j^2)^{k/2}$ gives $A_k r_j^{k+1}$. The squared norm of that neuron's three derivative columns is bounded pointwise by

$$t_k^2 2^k [k^2 c_j^2 (\gamma_j^2 + b_j^2)^{k-1} + (\gamma_j^2 + b_j^2)^k] \leq C_k^2 r_j^{2k}.$$

The maximum occurs at $(\gamma_j^2 + b_j^2)/r_j^2 = k/(k+1)$. Sum output bounds and squared column bounds, integrate against the probability measure, and use contractivity of P_H . With the Hilbert–Schmidt norm bounding the operator norm, this proves

$$\begin{aligned} \|q_H - P_H(\Psi_3 + \Psi_5)\| &\leq A_7 S_8, \\ \|J_H - J_3 - J_5\|_{\text{HS}} &\leq C_7 \sqrt{S_{14}}, \\ \|J_H^* f_\perp - G_3 - G_5\|_2 &\leq C_7 \|f_\perp\| \sqrt{S_{14}}. \end{aligned} \tag{A.7}$$

The same argument gives $\|q_H\| \leq A_3 S_4$ and $\|J_H\|_{\text{HS}} \leq C_3 \sqrt{S_6}$. These are global inequalities for the exact $\tanh$ network: no neuronwise small-argument hypothesis or polynomial replacement of training is used.

A.3 STRUCTURAL COEFFICIENTS AND TARGET DEPENDENCE

For $s > 0$ define

$$\begin{aligned} \chi_k &= W^{k/2-1} S_k / s^k, \\ \alpha_k &= W^{(k-1)/2} \|P_H \Psi_k\| / s^{k+1}, \\ \beta_k &= W^{(k-1)/2} \|J_k\|_{\text{HS}} / s^k, \quad \zeta_k = W^{(k-1)/2} \|G_k\|_2 / s^k, \quad k = 3, 5. \end{aligned} \tag{A.8}$$

These coefficients describe concentration, polynomial output, sensitivity, and target coupling. They are invariant under common positive rescaling of ϑ ; bounding them does not itself bound s . For equal neuron norms, $\chi_k = 1$; for the moments $k > 2$ used here, concentration in fewer neurons increases it. Let $\tau_k = \|P_{\mathcal{P}_k} f_\perp\|$, where $\mathcal{P}_k$ is the space of polynomials of degree at most k for the same measure. Then

$$\zeta_k \leq \beta_k \tau_k, \quad \alpha_k \leq A_k \chi_{k+1}, \quad \beta_k \leq C_k \sqrt{\chi_{2k}}, \quad k = 3, 5. \tag{A.9}$$

Each column of J_k is a polynomial of degree at most k , so $J_k^* f_\perp = J_k^* P_{\mathcal{P}_k} f_\perp$. The other inequalities follow from the fixed-radius bounds just proved. Thus cubic target orthogonality makes $G_3 = 0$, and quintic orthogonality also makes $G_5 = 0$. The mixed-sine and localized targets generally have nonzero low-degree coupling. The numerical evaluations retain $\alpha_k, \beta_k, \zeta_k$; substituting their concentration-only bounds gives a different, potentially looser evaluation, whose tightness is not claimed here.

At each GD iterate n , take finite nonnegative upper bounds $\hat{\alpha}_{k,n}, \hat{\beta}_{k,n}, \hat{\zeta}_{k,n}, \hat{\chi}_{8,n}, \hat{\chi}_{14,n}$ on these coefficients, and $\hat{r}_n \geq \|R_n\|_2$. Exact coefficients are allowed. For a candidate radius $B \geq 0$, define

$$\begin{aligned}\mathcal{Q}_n(B) &= \hat{\alpha}_{3,n} \frac{B^4}{W} + \hat{\alpha}_{5,n} \frac{B^6}{W^2} + A_7 \hat{\chi}_{8,n} \frac{B^8}{W^3}, \\ \mathcal{J}_n(B) &= \hat{\beta}_{3,n} \frac{B^3}{W} + \hat{\beta}_{5,n} \frac{B^5}{W^2} + C_7 \sqrt{\hat{\chi}_{14,n}} \frac{B^7}{W^3}, \\ \mathcal{T}_n(B) &= \hat{\zeta}_{3,n} \frac{B^3}{W} + \hat{\zeta}_{5,n} \frac{B^5}{W^2} + C_7 \|f_\perp\| \sqrt{\hat{\chi}_{14,n}} \frac{B^7}{W^3}, \\ \mathcal{V}_n(B) &= \mathcal{J}_n(B) \mathcal{Q}_n(B) + \mathcal{T}_n(B).\end{aligned}\tag{A.10}$$

For $s_n \leq B$, these bound $\|q_{H,n}\|$, $\|J_{H,n}\|$, $\|J_{H,n}^* f_\perp\|_2$, and $\|F_n\|_2$, respectively. For the last,

$$\|F_n\|_2 \leq \|J_{H,n}\| \|q_{H,n}\| + \|J_{H,n}^* f_\perp\|_2 \leq \mathcal{V}_n(B).\tag{A.11}$$

All four polynomials are nondecreasing in B . Their coefficients may evolve; neither the Jacobian nor the target coupling is frozen.

A.4 PRECISE GD THEOREM AND PROOF

Theorem A.1 (Population growth and output error for joint GD). *Consider (A.1) from an arbitrary restart, indexed by $n = 0$. Suppose $K_{C,n}$ is invertible, $s_n > 0$, and the structural upper bounds above hold for every iterate $0 \leq n \leq N$. Define*

$$B_0 = s_0, \quad B_{n+1} = B_n + \eta_n [\mathcal{V}_n(B_n) + \hat{r}_n].\tag{A.12}$$

Whenever these quantities are finite, for $0 \leq n \leq N$,

$$s_n \leq B_n, \quad \sum_{k < n} \|\vartheta_{k+1} - \vartheta_k\|_2 \leq B_n - B_0,\tag{A.13}$$

$$\frac{\|f - q_n\|}{\|f\|} \geq \frac{[\|f_\perp\| - \mathcal{Q}_n(B_n)]_+}{\|f\|}, \quad \bar{\lambda}_n \leq \lambda_{\text{RMS},n} \leq \frac{h B_n}{\sqrt{W}}.\tag{A.14}$$

Here $h > 0$ is the reference spacing, $\lambda_{\text{RMS},n} = h \|\gamma_n\|_2 / \sqrt{W}$, and $\bar{\lambda}_n = h W^{-1} \sum_j |\gamma_{j,n}|$.

Proof. Initially $s_0 = B_0$. If $s_n \leq B_n$, the exact GD update and (A.11) give

$$\|\vartheta_{n+1} - \vartheta_n\|_2 \leq \eta_n \|F_n + R_n\|_2 \leq \eta_n [\mathcal{V}_n(B_n) + \hat{r}_n] = B_{n+1} - B_n.$$

The triangle inequality implies $s_{n+1} \leq B_{n+1}$; induction and summation prove (A.13). Orthogonal projection and the reverse triangle inequality yield

$$\|q_n - f\| \geq \|q_{H,n} - f_\perp\| \geq [\|f_\perp\| - \|q_{H,n}\|]_+ \geq [\|f_\perp\| - \mathcal{Q}_n(B_n)]_+.$$

Finally Cauchy–Schwarz and $\|\gamma_n\|_2 \leq s_n$ prove both bandwidth bounds. These arguments use the discrete updates directly, with no continuous-time approximation or unmeasured finite-step correction. $\square$

The result holds for any nonnegative step sizes for which its premises hold. Large steps can make the comparison uninformative. Small post-transient tracking makes it useful but is not silently set to zero. A large value of B_n or a zero output floor signals loss of this sufficient explanation; it does not establish successful acquisition. Coefficient bounds checked only at stored times do not certify their values at every intervening GD step. For a pooled mean across equal-sized seeds, the per-seed bounds average; the theorem does not interpret neuron spread as uncertainty across seeds.

B Evidence and interpretation for Section 3.5

B.1 FIVE-MILLION-UPDATE OBSERVATION AND FEATURE ASSAYS

The primary target is $[\sin(2\pi x) + \frac{1}{2}\sin(6\pi x) + \frac{1}{4}\sin(14\pi x)]/\sqrt{21/32}$ on $[-1, 1]$. Total width is 512, including 22 halo nodes per side; the reference spacing is $h = 2/467$. Full-batch FP64 training uses 2,048 midpoints and five paired initializations. One recipe per joint optimizer is selected by median final error on 4,096 validation midpoints, requiring finite endpoints for every seed. The 8,192-point evaluation grid is a previously inspected resolution check, not an untouched test set. The selected joint Adam and GD recipes use initial rates 0.05 and 0.2, respectively, with cosine decay over all five million updates. The supplied-feature Adam comparison uses $\lambda = 1/4$, zero readout initialization, and the selected full-horizon cosine rate 10^{-4} . Both constant and cosine candidates are retained in the study records.

Figure 4a extends the same target, sampling grids, seed convention, and update budget to total widths $W \in \{128, 256, 512, 1024\}$. The corresponding numbers of interior intervals are $N \in \{105, 225, 467, 961\}$, with $h_W = 2/N$ and $R \in \{11, 15, 22, 31\}$ halo nodes per side; thus $W = N + 1 + 2R$. The initialization is the affine Xavier convention used above, with its random-number generator keyed by seed and the budget-derived interval count. Each width tests Adam cosine rates $\{0.02, 0.05\}$, Adam constant rates $\{0.002, 0.02\}$, and GD rates $\{0.2, 0.5\}$ under both schedules. Cosine decay spans the full five-million-update horizon. Selected Adam uses cosine rate 0.02 at width 128 and 0.05 otherwise. Selected GD uses rate 0.2 throughout, constant at widths 128 and 256 and cosine at widths 512 and 1024. These selections use the final validation criterion above, with no seed-specific recipe choice.

In increasing width order, mean absolute slopes are 4.324, 4.993, 3.943, and 2.295 for Adam, and 2.305, 1.072, 0.3171, and 0.1580 for GD. The supplied $\lambda = 1/4$ construction has slopes 13.125, 28.125, 58.375, and 120.125. The means do not keep pace with this reference over the tested widths. This comparison is conditional on the initialization, candidate grid, and budget; it does not establish an asymptotic growth law or a universal necessary slope. The fixed h_W is a reference spacing, not the spacing between learned centers.

Median joint training errors are 1.8143227×10^{-3} for Adam and 0.24393137 for GD. Supplied-feature Adam reaches 6.6245807×10^{-7} . The corresponding dense-grid values are 1.8148250×10^{-3} , 0.24393487, and 6.6198225×10^{-7} . Joint Adam’s median endpoint error falls about 28% between the two- and five-million-update experiments. This is a precision gap within the stated budget, not permanent failure. The pooled endpoint mean bandwidths are 0.0168845 for Adam and 0.00135795 for GD. Their neuron population standard deviations are 0.0541204 and 0.00703476. The supplied bandwidth is a comparison geometry, not a target-independent necessary threshold. Figure 4c uses the mean over all $5W$ neurons and its population standard deviation (divisor $5W$), not the standard deviation of the five seed means. The displayed lower edge is clipped at zero. Parameter snapshots are stored every 10,000 updates. Panel (b) retains within-bin error extrema and exact endpoints; its shaded region is not a confidence interval.

All 160 seed/recipe trajectories reached the full horizon. An independent NumPy evaluation verified endpoint output errors. The selected width-512 Adam slopes and intercepts reproduce the five archived feature-assay dictionaries exactly, so the following interventions concern the same learned features.

Freezing the five-million-update Adam features and training new readouts for two million additional updates gives median dense-grid error 1.8147369×10^{-3} . SVD readout fits at relative cutoffs 10^{-12} and 10^{-14} also yield only small improvements; lowering the cutoff to 10^{-16} exposes sensitivity to floating-point cancellation. These assays support a limitation of the learned dictionary at the checked numerical precision, not an exact-arithmetic approximation floor. Multiplying both slope and intercept by 4 or 16 preserves each learned center. After the same readout-training assay, median dense-grid errors are 0.0916060 and 0.138935, respectively. Thus simple slope amplification does not repair this geometry. These interventions change nonnested feature spaces; they do not establish that

larger slopes always help. The assay’s additional optimization budget is separate from the end-to-end five-million-update comparison.

B.2 CONDITIONAL THEOREM EVALUATION ON EVOLVING GD POPULATIONS

The mechanism study uses a separate cohort: total width 705, reference spacing $h = 1/256$, 2,048 uniform midpoint samples, and restarts after 20,000 ordinary-GD updates at step size 0.002. Each continuation runs another 100,000 updates, or 200 units of flow time. Seed 30 covers six targets: an orthogonal degree-five mixture, mixed sine, a shifted Gaussian, a compact bump, a smooth step, and an absolute-value kink. The mixture is $0.3P_0 + 0.4P_1 + \sqrt{0.75}P_5$, with P_k orthonormal for the empirical measure. The other unnormalized targets are

$$\begin{aligned} & \sin(2\pi x) + \frac{1}{2} \sin(6\pi x) + \frac{1}{4} \sin(14\pi x), \\ & e^{-((x+0.35)/0.22)^2}, \quad \tanh(14(x - 0.31)), \quad |x + 0.23|, \\ & f_{\text{bump}}(x) = \begin{cases} \exp\left(1 - \frac{1}{1-u^2}\right), & |u| < 1, \\ 0, & |u| \geq 1. \end{cases} \quad u = (x - 0.35)/0.22, \end{aligned}$$

Each is divided by its RMS on the original training grid. Four additional seed-31 restarts cover step, kink, Gaussian, and bump; they were selected for weak initial comparison margins before their continuation outcomes were inspected.

The discrete scalar recurrence uses structural diagnostics sampled every unit of flow time, interpolating its nonnegative polynomial coefficients and tracking norm to native GD steps. This implements a numerical evaluation of the conditional theorem; interpolation is not a proved upper enclosure between stored samples. In the six principal GD runs the minimum sampled relative-error floors are, in the target order above, 0.865734, 0.912194, 0.784451, 0.879212, 0.479360, and 0.426952. The corresponding executed endpoint errors are 0.866025, 0.912564, 0.788073, 0.881027, 0.489675, and 0.427513. The minimum floors and endpoint errors are distinct summaries, not an every-step tightness statistic. All ten growing comparisons remain useful over the full continuation, with seed-31 minimum floors from 0.426695 to 0.878467. Principal endpoint RMS-bandwidth bounds range from 3.63×10^{-4} to 5.18×10^{-4} , versus observed values from 2.09×10^{-4} to 2.51×10^{-4} .




Figure 6: **The output-error bound remains informative across all ten continuations.** The theorem’s lower bound divided by executed relative training RMS error at each stored diagnostic time, every 500 GD updates. Solid curves are seed 30; dashed curves are the four seed-31 stress cases. The smallest ratio is 0.97714. These are evaluations with interpolated structural coefficients from the trajectories, not initial-state forecasts or certified bounds between diagnostic samples.

Figure 6 compares bound and error at matching times throughout the trajectories. The relative shortfall never exceeds 2.29% at stored diagnostic times. This statistic directly measures output-bound tightness, separately from the endpoint summaries above.

The saved verification includes independent automatic derivatives, global remainder checks, exact discrete moment identities, and a solvable positive-growth comparison. Step refinement changes population size and normalized slope RMS by less than 4.8×10^{-7} relatively across both seeds. Decimating the structural diagnostics changes the final GD comparison radius by less than 9×10^{-6} relatively. These checks support numerical resolution; they do not supply interval-arithmetic certificates for the conditional premises.

The theorem explains this post-transient regime through evolving population structure. The older width-177 archive includes regimes where large moments make the polynomial comparison uninformative. No coverage claim for all 23 archived targets, all training ages, or the full five-million-update Figure 4 trajectories follows. Adam’s coordinate-dependent scaling and momentum also fall outside the ordinary-GD proof. Its observed precision gap motivates the same empirical question without establishing the same quantitative mechanism.